

# Toward a Clearer Characterization of Neuro-Symbolic Frameworks: A Brief Comparative Analysis

Sania Sinha

Tanawan Premisri

Parisa Kordjamshidi

*Department of Computer Science and Engineering, Michigan State University*

SINHASA3@MSU.EDU

PREMSRIT@MSU.EDU

KORDJAMS@MSU.EDU

**Editors:** Leilani H. Gilpin, Eleonora Giunchiglia, Pascal Hitzler, and Emile van Krieken

## Abstract

Neurosymbolic (NeSy) frameworks combine neural representations and learning with symbolic representations and reasoning. Combining the reasoning capacities, explainability, and interpretability of symbolic processing with the flexibility and power of neural computing allows us to solve complex problems with more reliability while being data-efficient. However, this recently growing topic poses a challenge to developers with its learning curve, lack of user-friendly tools, libraries, and unifying frameworks. In this paper, we characterize the technical facets of existing NeSy frameworks, such as the symbolic representation language, integration with neural models, and the underlying algorithms. A majority of the NeSy research focuses on algorithms instead of providing generic frameworks for declarative problem specification to leverage problem solving. To highlight the key aspects of Neurosymbolic modeling, we showcase three generic NeSy frameworks - *DeepProbLog*, *Scallop*, and *DomiKnowS*. We identify the challenges within each facet that lay the foundation for identifying the expressivity of each framework in solving a variety of problems. Building on this foundation, we aim to spark transformative action and encourage the community to rethink this problem in novel ways.

**Keywords:** Neurosymbolic, Comparing NeSy frameworks, DomiKnowS, DeepProbLog, Scallop, Combining learning and reasoning

## 1. Introduction

**Symbolic or good old-fashioned AI** focused on creating rule-based reasoning systems (Hayes-Roth, 1985) exemplified with early works such as the Physical Symbol System (Augusto, 2021; Newell, 1980) and ELIZA (Weizenbaum, 1966). However, drawbacks such as limited scalability due to the need to explicitly define domain predicates and rules for each task, lack of robustness in handling messy real-world data, and low computational efficiency led to a decline in the popularity of this paradigm, shifting the focus toward neural computing and deep learning. **Deep Learning** (LeCun et al., 2015; Ahmad et al., 2019) revolutionized AI as nuanced relationships in data could be learned by backpropagation through multiple layers of processing and creating abstract representations of data. However, it led to a loss of explainability (Li et al., 2023a), dependence on large amounts of data, and rising concerns about its environmental sustainability (Bender et al., 2021). **Neurosymbolic AI** (Hitzler and Sarker, 2022; Bhuyan et al., 2024), a combination of symbolic AI and reasoning with neural networks, attempts to incorporate the capabilities of both worlds and create systems that are data and time efficient, generalizable, and explainable.

Neurosymbolic models have been applied to several real-world applications (Bouneffouf and Aggarwal, 2022) in safety-critical areas (Lu et al., 2024) such as healthcare (Hossain and Chen, 2025) and autonomous driving (Sun et al., 2021). Several techniques have been proposed for this integration (Kautz, 2022; Jayasingha et al., 2025), trying to combine the pros and mitigate the cons from both symbolic and neural methods. However, due to lack of unified libraries to facilitate this research and the focus on specific algorithms rather than generic frameworks, this research becomes less impactful. Moreover, the few generic frameworks tend to vary in problem formulation, implementation, algorithms, and flexibility of use. This poses a challenge in being able to compare their performance uniformly or identify a research direction that improves on previous work. To alleviate this issue, we provide a comparative study with the following key contributions.

a) Identifying the main components of existing NeSy frameworks, b) Comparison of frameworks across the identified facets, c) Highlighting the requirements for the next generation of NeSy frameworks, building upon the drawbacks of the current systems and the possible interplays between the neural and symbolic components. We plan to expand this study to cover more frameworks while the three selected ones are used to explain the aspects of our characterization. These frameworks are demonstrated with an example task detailed in Appendix A, tying the comparative facets concretely with a technical implementation.<sup>1</sup> The MNIST Sum is a modified version of the classic MNIST digit recognition task (Lecun et al., 1998), where a model is given images of two digits and asked to predict their sum.

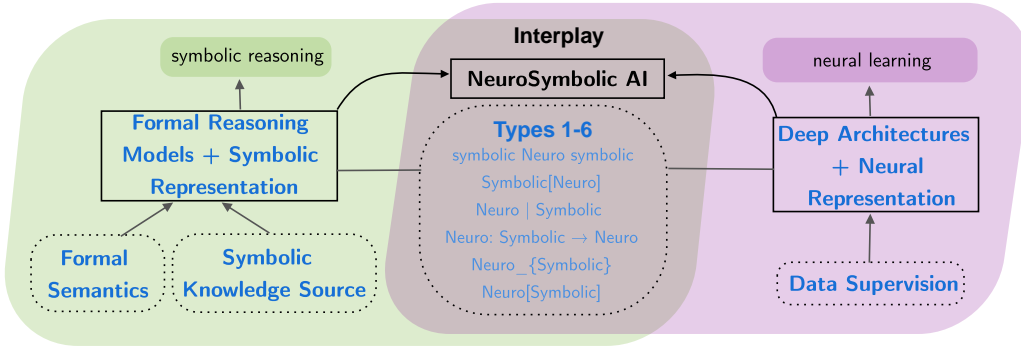


Figure 1: An overview of the main components of a neurosymbolic framework.

## 2. Neurosymbolic Frameworks

A NeSy framework should provide flexibility for modeling both neural and symbolic components (Kordjamshidi et al., 2016, 2015) and their interplay in a unified declarative framework, going beyond specific underlying algorithms and techniques. On the symbolic side, a generic framework should support a symbolic representation language that can be seamlessly connected to neural components and cover different symbolic reasoning mechanisms. On the neural side, we need to have the flexibility of connecting to various architectures, including various loss functions, sources of supervision, and training paradigms. More importantly, a

1. <https://github.com/HLR/nesy-examples>

Framework	Symbolic		Model Dec	Interplay		LLM
	Lang	Knowledge Rep		Algorithm	Eff	
CCN+	None	Propositional Logic Clauses	✗	ReqL & ReqLoss	✗	✗
DomiKnowS	None	Concepts, Constraints	✓	Primal-Dual, Sampling Loss	✗	<a href="#">Faghihi et al. (2024)</a>
DeepProbLog	ProbLog	Facts, Rules, Predicates	✗	Entailment	✗	✗
LEFT	None	First Order Logic	✗	Differentiable Reasoning	✗	<a href="#">Hsu et al. (2023)</a>
PyReason	None	Constants, Relations, Facts, Rules	✗	Reasoning over graph	✓	✗
Scallop	DataLog	Rules, Relations	✗	Differentiable Reasoning	✓	<a href="#">Li et al. (2024)</a>

Table 1: Frameworks with their comparative factors. Lang: External language required, Knowledge Rep: Knowledge Representation, Model Dec: Model Declaration flexibility, Algorithm: Supported algorithm(s) for learning and inference, Eff: Computational efficiency considerations, LLM: Use of Large Language Models.

NeSy framework should provide a modeling language for specification and seamless integration of the two components in building pipelines or arbitrary composition of models. Such a NeSy framework should support neuro-symbolic training and inference beyond specific integration algorithms. We distinguish between *NeSy techniques* and *NeSy frameworks*. By techniques, we mean when task-specific solutions are provided ([Lample and Charton, 2020](#); [Burattini et al., 2002](#)). For example, **AlphaGo** ([Silver et al., 2016](#)) introduced a reinforcement learning solution to Go, using Monte Carlo Tree Search as a symbolic component inside a neural network. Another example is **NS-CL** ([Mao et al., 2019](#)) (Neuro-Symbolic Concept Learner) that integrates neural perception with symbolic reasoning to learn visual concepts and compositional language grounding for VQA tasks. Many other techniques and algorithms are proposed for the interplay between the two paradigms ([Badreddine et al., 2022](#); [Cohen et al., 2017](#); [Smolensky et al., 2016](#); [Lima et al., 2005](#); [Sathasivam, 2011](#); [Serafini and d’Avila Garcez, 2016](#); [Lamb et al., 2021](#)) such as Inference Masked Loss ([Guo et al., 2020](#)), Semantic Loss ([Xu et al., 2018](#)), Primal-Dual ([Nandwani et al., 2019](#)), etc., later discussed in Section 6. They often lack the generality of frameworks, which are designed as broader tools intended for practical use and extensibility with new integration algorithms and with the capability of programming and configuring the two parts and their interplay.

In this work, we focus on a selection of generic NeSy frameworks. The following are examples of research efforts towards advancing the development of such general-purpose frameworks: **DeepProbLog** ([Manhaeve et al., 2021](#)) is a probabilistic logic programming language, incorporating neural predicates in logic programming with an underlying differentiable translation of logical reasoning. The probabilistic logic programming component is built on top of ProbLog ([De Raedt et al., 2007](#)). **DomiKnowS** ([Rajaby Faghihi et al., 2021](#); [Faghihi et al., 2023, 2024](#)) is a declarative learning-based programming framework ([Kord-](#)

jamshidi et al., 2019) that integrates symbolic domain knowledge into deep learning. It is a Python framework, facilitating the incorporation of logical constraints that represent domain knowledge with neural learning in PyTorch. **Scallop** (Huang et al., 2021; Li et al., 2023b, 2024) is a framework that includes flexible symbolic representation based on relational data modeling, using a declarative logic programming built on top of Datalog (Abiteboul et al., 1995) with a framework for automatic differentiable reasoning. **LEFT** (Hsu et al., 2023) is a less generic framework designed for grounding language in visual modality and compositional reasoning over concepts. The framework consists of an LLM interpreter that converts natural language to logical programs. The generated programs are directed to a differentiable, domain-independent, and soft first-order logic-based executor. LEFT is limited to tasks requiring grounding language in vision such as visual question answering (Johnson et al., 2017; Yi et al., 2018; Liu et al., 2019). Building on this foundation, NeSyCoCo (Kamali et al., 2025) was introduced to address its limitations, particularly its struggle with lexical variety and handling unseen concepts. NeSyCoCo extends LEFT’s approach by using distributed word representations to connect a variety of linguistically motivated predicates to neural modules, to alleviate reliance on a predefined predicate vocabulary. **PyReason** (Aditya et al., 2023) is a library built to support reasoning on top of outputs from neural networks. The neural component produces outputs such as labels or concept scores. While the symbolic component does graph-based reasoning using logic rules declared over a graph structure. It can produce an explanation trace for inference and has a memory-efficient implementation. **PLoT** (Wong et al., 2023) (Probabilistic Language of Thought) is a *proposed* framework leveraging neural and probabilistic modeling for generative world modeling. It models thinking with probabilistic programs and meaning construction with neural programs. The goal is to provide a language-driven unified thinking interface. **CCN+** (Giunchiglia et al., 2024) is a framework that modifies the output layer of a neural network to make results compliant with requirements that can be expressed in propositional logic. A requirement layer, ReqL, is built on top of the neural network. The standard cross-entropy loss is adapted into a ReqLoss to learn from the constraints in the ReqL layer. We characterize frameworks by: a) Symbolic knowledge representation language, b) Representation and flexibility of Neural Modeling, c) Model Declaration, d) Interplay between symbolic and sub-symbolic systems, and e) Usage of LLMs. Figure 1 shows the relationship between these aspects. The neural representations and the symbolic representations are the two main components of a neurosymbolic framework. The neural representation guides learning and obtaining supervision from the data, while the symbolic representations leverage symbolic reasoning, where the symbolic knowledge can be exploited during training or inference. Table 1 shows an overview of the frameworks across chosen features. For future sections, we focus on **DomiKnowS**, **DeepProbLog**, and **Scallop** to provide a deeper investigation of the challenges in each component. Due to differences in implementation, each framework allows for easy implementation of different tasks. With the chosen frameworks, we can solve the same task across all.

### 3. Symbolic Knowledge Representation

Generic Neuro-Symbolic (NeSy) systems and frameworks use symbolic knowledge representation languages to encode constraints, facts, probabilities, and rules. Frameworks vary in how they represent and integrate this symbolic knowledge. Many employ classical formal

logic, grounded in well-defined syntax and semantics, and adapt these representations and reasoning mechanisms within a unified integration framework. Some frameworks build on established formalisms such as logic programming or constraint satisfaction. In contrast, others take an entirely new hybrid semantics, while preserving conventional symbolic syntax. Figure 2 compares the implementation of symbolic knowledge (concepts or facts) for the MNIST Sum task. In general, the domain knowledge consists of the two concepts of *digits* and the *sum*. As can be seen, Domiknows represents a part of symbolic domain knowledge as a graph  $G(V, E)$ , where the nodes are the concepts in the domain and the edges are the relationships between them. Each node can have properties. More complex knowledge beyond entities and relations is expressed with a pseudo first-order logical language with quantifiers designed in Python. DomiKnowS mostly interprets the symbolic knowledge as logical constraints, such as the implementation of `sum_combinations` in the given example. Unlike the other frameworks, DomiKnowS does not build on predefined formal semantics. It follows a FOL-like syntax for symbolic logical representations, making it independent of the formal semantics of an underlying formal language and allows more flexibility of representations and adaptation to underlying algorithms in the framework. DeepProbLog, on the other hand, utilizes logical predicates that are originally a part of the probabilistic logic programs (Ng and Subrahmanian, 1992) of ProbLog (De Raedt et al., 2007), for its symbolic representation. These neural predicates obtain their probability distributions from the underlying neural models. Probabilistic facts, neural facts, and neural annotated disjunctions (nAD) whose probabilities are supplied by the neural component of the program can be added. Here, `digit` is a neural predicate as indicated by the use of `nn(...)`. DeepProbLog follows the formal semantics of Prolog (Clocksin and Mellish, 2003), followed by ProbLog, its probabilistic extension. Scallop adopts a relational data model for symbolic knowledge representation (Kolaitis and Vardi, 1990). It is built on the syntax and formal semantics of Datalog and its probabilistic extensions, relaxing the exact semantics of ProbLog. It allows for the expression of common reasoning, such as aggregation, negation, and recursion. Similar to DeepProbLog, some of these predicates in the symbolic part obtain their probability distribution from neural models, such as `digit_1` and `digit_2`. While ProbLog requires exhaustive search for computations, Datalog can use top-k results and exploit database optimizations, making Scallop algorithmically more time-efficient.

Here, `digit` is a neural predicate as indicated by the use of `nn(...)`. DeepProbLog follows the formal semantics of Prolog (Clocksin and Mellish, 2003), followed by ProbLog, its probabilistic extension. Scallop adopts a relational data model for symbolic knowledge representation (Kolaitis and Vardi, 1990). It is built on the syntax and formal semantics of Datalog and its probabilistic extensions, relaxing the exact semantics of ProbLog. It allows for the expression of common reasoning, such as aggregation, negation, and recursion. Similar to DeepProbLog, some of these predicates in the symbolic part obtain their probability distribution from neural models, such as `digit_1` and `digit_2`. While ProbLog requires exhaustive search for computations, Datalog can use top-k results and exploit database optimizations, making Scallop algorithmically more time-efficient.

#### 4. Neural Models Representations

The other core component of a NeSy system is the neural modeling that is integrated with the symbolic knowledge discussed above. The neural models are mostly wrapped up under the logical predicate names in most of the frameworks that have an explicit

Domiknows	DeepProbLog
<pre>image = Concept() digit = image(ConceptClass=     NumericalConcept)  image_pair = Concept() (pair_d0,  pair_d1) = image_pair.has_a(     digit0=image,     digit1=image)  s = image_pair(     ConceptClass=EnumConcept,     values=summations)  ....  sum_combinations.append(     andl(         getattr(digit, d0_nm)         (path='x', pair_d0)),         getattr(digit, d1_nm)         (path='x', pair_d1))))  ....</pre>	<pre>nn(m_digit, [X], Y,     [0....9]):: digit(X,Y).  addition(X,Y,Z) :-     digit(X,X2), digit(Y,Y2),     Z is X2+Y2.</pre>
Scallop	
<pre>scl_ctx.add_relation(     "digit_1", int,     input_mapping=         list(range(10))) self.scl_ctx.add_relation(     "digit_2", int,     input_mapping=         list(range(10))) self.scl_ctx.add_rule(     "sum_2(a + b)     :- digit_1(a), digit_2(b)") self.sum_2 =     self.scl_ctx.forward_function     ("sum_2", output_mapping=         [(i,) for i in range(19)])  ....</pre>	

Figure 2: Comparison of Symbolic Representation across frameworks.

logical knowledge representation language. To leverage the reasoning capabilities of the symbolic system available and the ability of neural models to learn abstract representations from data, the neural models are used as abstract concept learners for the concepts defined as logical predicates in the symbolic representation. The neural model representation is often used to predict probability distributions for the symbolic concepts based on raw sensory inputs. The neural modeling is often written using standard deep learning libraries, such as PyTorch (Paszke et al., 2019). Figure 3 shows snippets of neural modeling expressions across frameworks, highlighting differences in implementation. Scallop utilizes relatively standard neural modeling using PyTorch, while needing an added context of symbolic rules. Although integrated into Python, the context relation and rule setup are verbatim from DataLog and only passed as a parameter to a function, which requires familiarity with DataLog and its semantics. DeepProbLog, on the other hand, needs manual configuration of the raw data and processing into queries, on top of other standard neural components. This processed data is passed into the neural network which is then connected to a ProbLog program, such as `addition.pl` in the figure. DomiKnowS’s neural component is built in PyTorch. Unlike other frameworks, DomiKnowS has built-in components called *Readers*, *Sensors* and *Module learners* that make the connection to neural components and feeding data to them explicit in the program. This provides more flexibility in connecting the concepts to deterministic or probabilistic functions that can interact with other symbolic concepts. The module learner can also use custom models. This makes the interaction with raw data structured, transparent, and controllable.

DomiKnowS	DeepProbLog
<pre> class Net:     # neural network  image['pixels'] = ReaderSensor (keyword='pixels') image_batch['pixels'], image_contains.reversed]= JointSensor(image['pixels'], forward=make_batch) image['logits'] = ModuleLearner( 'pixels', module=Net()) ... </pre>	<pre> class MNIST_Net:     # neural network  network = MNIST_Net() net = Network(network, "mnist_net", batching=True) net.optimizer = torch.optim.Adam (network.parameters(), lr=1e-3)  model = Model("models/addition.pl", [net]) ... </pre>
Scallop	
<pre> class MNISTSum2Net(nn.Module):     def __init__(self, provenance, k):         self.mnist_net = MNISTNet()         self.scl_ctx = scallop.ScallopContext(provenance=provenance, k=k)     ...  class Trainer():     def __init__(self, train_loader, test_loader, model_dir, learning_rate, loss, k, provenance):         self.model_dir = model_dir         self.network = MNISTSum2Net(provenance, k)         self.optimizer = optim.Adam(self.network.parameters(), lr=learning_rate)     ... </pre>	

Figure 3: For neural integration, DomiKnowS employs sensors and readers for reading in data, while a learner connects to a network. DeepProbLog connects the neural network to the ProbLog file, requiring data handling to construct the terms and queries from raw data. Scallop has an extra layer on top of the standard network that adds the symbolic context.

## 5. Model Declaration

Most frameworks utilize neural components as abstract concept learners and use a symbolic component to reason over the learned concepts. Each learner is a model and *model declaration* refers to the flexibility of modularizing and connecting different learners. Each learner can receive supervision independently. In most neurosymbolic frameworks, the supervision from data is usually provided based on the final output of the end-to-end model. For example, in an MNIST Sum task used throughout and detailed in Appendix A, the neural and symbolic components are trained based on the final output of the sum, without access



to individual digit labels in a semi-supervised setting. The task loss, e.g., a Cross-Entropy Loss, is computed, and errors are backpropagated through the differentiable operations that led to the output generation. For example, in DeepProbLog, we can declare a single loss function associated with the entire neural component. Gradient computations differ across frameworks depending on whether losses are defined individually for each neural output or specified as a single global loss function. However, there remains a need for models capable of incorporating supervision at multiple levels of their symbolic representations. In DomiKnowS, loss computation can be defined for each symbol. Since each concept is linked to both learning modules and ground-truth labels, their losses can be integrated seamlessly. This enables joint training of all concepts alongside the target task, allowing each concept to be optimized more effectively—leveraging available data without relying solely on the target task’s output. In other words, it provides the flexibility of building pipelines of decision making, obtaining distant supervision in addition to joint training and inference.

## 6. Interplay between Symbolic and Sub-symbolic

Kautz (2022) provides a characterization of the possible interplays between symbolic and sub-symbolic components. This interplay of neurosymbolic can be explained by the concept of System 1 and System 2 thinking described in Kahneman (2011). Research in this field aims to create an ideal integration that seamlessly supports ”thinking fast and slow” (Booch et al., 2021; Fabiano et al., 2023). Here, System 1 refers to the fast neural processing, while System 2 corresponds to the slower, more deliberate symbolic reasoning. *Different methods for the integration of symbolic reasoning and neural programming have been explored such as employing logical constraint satisfaction, integer linear programming, differentiable reasoning, probabilistic logic programming.* In this section, we will discuss a system-level algorithmic comparison of the different frameworks.

**DomiKnowS** models the inference as an integer linear programming problem to enforce the model to follow constraints expressed in first-order logical form (Van Hentenryck et al., 1992). The objective of the program is guided by the neural components, and the framework supports multiple training algorithms for learning from constraints. The Primal-Dual formulation (Nandwani et al., 2019) converts the constrained optimization problem into a min-max optimization with Lagrangian multipliers for each constraint, augmenting the original loss with a soft logic surrogate to minimize constraint violations. Sampling-Loss (Ahmed et al., 2022), inspired by semantic loss (Xu et al., 2018) and samples a set of assignments for each variable based on the probability distribution of the neural modules’ output. Integer Linear Programming (ILP) (Cropper and Dumančić, 2022) formulates an optimization objective based on Inference-Masked Loss (Guo et al., 2020) to constrain the model during training. The training goal is to adjust the neural models to produce legitimate outputs that adhere to the given constraints. At prediction time, ILP can also be applied to enforce final predictions that comply with given constraints. DomiKnowS relies on the off-the-shelf optimization solver Gurobi (Gurobi Optimization, LLC, 2024). **DeepProbLog** models each problem as a probabilistic logical program that consists of neural facts, probabilistic facts, neural predicates, and a set of logical rules. A joint optimization of the parameters of the logic program is done alongside the parameters of the neural component. Neural network training is done using learning from entailment (Frazier and Pitt, 1993) while in ProbLog, gradient-based optimization is performed on the underlying

generated Arithmetic Circuits (Shpilka et al., 2010), which is a differentiable structure. The Arithmetic Circuits are transformed from a Sentential Decision Diagram (Darwiche, 2011) generated by ProbLog. Algebraic ProbLog (Kimmig et al., 2011) is used to compute the gradient alongside probabilities using semirings (Eisner, 2002). **Scallop** is similar in its setup to DeepProbLog where it creates an end-to-end differentiable framework combining a symbolic reasoning component with a neural modeling component. They aim to relax the formal semantics required by the use of ProbLog in DeepProbLog and instead rely on a symbolic reasoning language extending DataLog, built into their framework. They have a customizable provenance semiring framework (Green et al., 2007), where different provenance semirings, such as extended max-min semiring and top-k proofs semiring, allow learning using different types of heuristics for gradient calculations. Table 2 compares the computational efficiency of these models at training and inference time on a single training/testing example. As theoretically suggested, Scallop is expected to outperform other frameworks in inference and training speed, owing to its memory and time-efficient implementation in Rust. The results in Table 2 support this expectation, with Scallop achieving the fastest inference time, on par with DomiKnowS. In practice, DeepProbLog achieves slightly faster training performance than Scallop. This discrepancy may be due to overhead unrelated to the core algorithmic complexity. DomiKnowS exhibits slower training due to the overhead of uploading the entire graph of data into memory.

Framework	Training Time(ms)	Inference Time(ms)
DomiKnowS	37.72	<b>2.34</b>
DeepProbLog	<b>5.84</b>	3.24
Scallop	6.50	<b>2.35</b>

Table 2: Computation time in milliseconds (ms) for training and inference across frameworks based on one training/testing example.

## 7. Role of Large Language Models

Large foundation models hold significant promise for overcoming the bottleneck of acquiring symbolic representations, which are essential for symbolic reasoning and consequently in neurosymbolic frameworks. **Source of Symbolic Knowledge:** The symbolic knowledge in neuro-symbolic systems, which is integrated with the neural component, can originate from several distinct sources. While most systems require explicit, hand-crafted symbolic knowledge, earlier classical logic-based learning research can be used for automatically learning rules from data by using inductive logic programming (Nienhuys-Cheng and de Wolf, 1997; Bratko and Muggleton, 1995) or mining constraints. Nowadays, even LLMs can be utilized to generate symbolic knowledge (Pan et al., 2023a; Mirzaee and Kordjamshidi, 2023; Acharya et al., 2024; Xu et al., 2024a). Several neurosymbolic frameworks and systems have tried utilizing large foundation models to generate the symbolic knowledge, based on the task or query, to overcome the labor-intensive nature of hand-crafting rules for every single task and the time required in the automatic learning of symbolic knowledge from data (Ishay et al., 2023; Xu et al., 2024a; Yang et al., 2024). Extraction of symbolic representations from Foundation Models has become possible given the vast implicit knowledge stored within these models, such as LLMs and multimodal models, which are trained on massive and diverse corpora (Li et al., 2024; Petroni et al., 2019). These models can gen-



erate symbolic content (e.g., candidate rules, knowledge graph triples, or logic statements), perform reasoning that mimics symbolic inference, or act as components alongside symbolic modules (Fang and Yu, 2024). For example, LLMs can be prompted to extract facts from unstructured text, effectively populating a symbolic knowledge graph (Yao et al., 2025). Techniques like Symbolic Chain-of-Thought inject formal logic into the LLM’s reasoning process, improving accuracy and explainability on logical reasoning tasks (Xu et al., 2024b). However, foundation models are prone to hallucinations and lack the strict logical guarantees of traditional symbolic systems (Zheng et al., 2024). Therefore, integrating foundation models often requires careful prompting, verification steps to ensure reliability (Xu et al., 2024b). **Generation of inputs to symbolic engines:** LLMs have also been used to generate translations from raw inputs, specially natural language, to symbolic language that is then fed into a symbolic reasoner. For example in Logic-LM (Pan et al., 2023b), LLMs are leveraged to convert a natural language query into symbolic language that is then solved by a symbolic reasoner. This method improves the performance of unfinetuned LLMs on logical reasoning-based tasks. DomiKnowS (Faghihi et al., 2024) takes this a step further by enabling users to describe problems in natural language which LLMs then use to generate relevant concepts and relationships. Through a user-interactive process, these concepts and relationships are refined iteratively. Finally, the LLM translates the user-defined constraints from natural language into first-order logic representations before converting them into DomiKnowS syntax. Some systems use LLMs in multiple capacities. In VIERA (Li et al., 2024), which is built on top of Scallop, 12 foundation models can be used as plugins. These models are treated as stateless functions with relational inputs and outputs. These foundation models can be either language models like GPT (OpenAI et al., 2024) and LLaMA (Touvron et al., 2023), vision models such as OWL-ViT (Minderer et al., 2022) and SAM (Kirillov et al., 2023), or multimodal models such as CLIP (Radford et al., 2021). These models can be used to extract facts, assign probabilities, or for classification, and are treated as "foreign predicates" in their interface. An older version, DSR-LM (Zhang et al., 2023) utilized BERT-based language models for perception and relation extraction, combined with a symbolic reasoner for question answering. LEFT uses LLMs both for generation of the concepts used for grounding and as an interpreter to generate the first-order logic program for a natural language query, that is solved by the symbolic executor.

## 8. Discussion and Future Direction

Table 1 summarizes the comparative aspects of existing frameworks and outlines future directions for optimizing, as we observe many columns marked with 'X', implying most frameworks present challenges that hinder the application and flexibility of the frameworks. While current frameworks are functional, future developments should take a more holistic approach that considers all aspects from an end-user perspective, aiming to improve usability as general-purpose libraries and foster wider adoption of neurosymbolic methods. **Symbolic Representation.** The generic neurosymbolic frameworks provide a formal knowledge representation language of their choice. The selected languages often are based on pure logical formalisms with established formal semantics, for example, Datalog or Prolog. However, we argue that knowledge representation for neurosymbolic frameworks needs to be an innovative language designed for this integration purpose with adaptable semantics with learning as the pivotal concept (Kordjamshidi et al., 2019). Restricting these frame-

works to classical AI formalisms and formal semantics limits the level of extension that can be made and restricts the support of various algorithms and types of integration.

**Neural Modeling.** Most of the examined frameworks leave neural modeling and the task of connecting the symbolic and sub-symbolic components to the user. This connection usually requires low-level data preprocessing, that’s time consuming to implement. A lack of user-friendly libraries discourages developers from using neurosymbolic methods to solve downstream tasks. These frameworks need abstractions (Kordjamshidi et al., 2022) that improve user experience and remove the need to implement data processing from scratch.

**Model Declaration.** There is a need to be explicit about the low-level components of the neural architecture, enabling us to design interactions between neural and symbolic components and connect them as intended. The goal is to provide flexibility in designing arbitrary loss functions and connecting them to data for supervising concepts at various neural layers, which will allow any symbol to be learnable.

**Types of Interplay.** Considering Kautz (2022)’s classification, current frameworks are limited in supporting one or two ways of interactions. The "Algo" column in Table 1 shows that DeepProbLog and Scallop utilize one form of implementation, while DomiKnowS has multiple settings. One of the key challenges is determining the appropriate level of abstraction in a neural model after which reasoning should occur. The classification types demonstrate how a neural model can identify the relevant symbolic representations and suggest that neurosymbolic frameworks could leverage these models to learn and route inputs to the corresponding symbolic reasoning system. However, it remains unclear what level of abstraction is most effective for solving the end task in practice.

**LLM.** Drawbacks often associated with employing symbolic AI into neural computing, such as creation of the symbolic knowledge for integration, can be mediated with the use of LLMs and foundation models. LLMs have the potential to alleviate the classical issues in symbolic processing. Their vast knowledge can also be utilized to reduce the need for rebuilding neural components, allowing for flexible connections with different symbolic components.

## 9. Conclusion

Neurosymbolic AI presents a promising path forward in addressing the limitations of purely symbolic or neural approaches to AI. By integrating symbolic reasoning with neural learning, NeSy frameworks offer a balance between interpretability, data and time efficiency, and generalization. In this paper, we characterize the core components of NeSy frameworks and provide an analysis of some existing ones - DeepProbLog, Scallop, and DomiKnowS, illustrating the comparative facets. We identified some facets as symbolic knowledge and data representation, neural modeling, model declaration, method of integrating the symbolic and sub-symbolic systems, and role of LLMs. We identify key challenges in each facet that can guide us toward building the next generation of neurosymbolic frameworks. Unifying ideas in the field and building flexible frameworks by incorporating strengths in every facet will ease the learning curve associated with NeSy systems and improve standardization. Future NeSy frameworks should aim to provide flexible implementation, a user-friendly interface, improve scalability, and develop seamless integrations with foundation models. The advent of next-generation LLMs/VLMs provides promising solutions to longstanding knowledge engineering challenges, fostering more effective and scalable integration of symbolic representations and advancing research in neurosymbolic AI.

## Acknowledgments

This project is partially supported by the Office of Naval Research (ONR) grant N00014-23-1-2417. Any opinions, findings, conclusions, or recommendations expressed in this material are those of the authors and do not necessarily reflect the views of the Office of Naval Research. We thank Danial Kamali for his help in early stages of drafting and Uzair Mohammad for his editing and suggestions for Figure 1.

## References

- Serge Abiteboul, Richard Hull, and Victor Vianu. *Foundations of Databases: The Logical Level*. Addison-Wesley Longman Publishing Co., Inc., USA, 1st edition, 1995. ISBN 0201537710.
- Kamal Acharya, Alvaro Velasquez, and Houbing Herbert Song. A survey on symbolic knowledge distillation of large language models. *IEEE Transactions on Artificial Intelligence*, 5(12):5928–5948, 2024. doi: 10.1109/TAI.2024.3428519.
- Dyuman Aditya, Kaustuv Mukherji, Srikar Balasubramanian, Abhiraj Chaudhary, and Paulo Shakarian. Pyreason: Software for open world temporal logic, 2023. URL <https://arxiv.org/abs/2302.13482>.
- Jamil Ahmad, Haleem Farman, and Zahoor Jan. *Deep Learning Methods and Applications*, pages 31–42. Springer Singapore, Singapore, 2019. ISBN 978-981-13-3459-7. doi: 10.1007/978-981-13-3459-7\_3. URL [https://doi.org/10.1007/978-981-13-3459-7\\_3](https://doi.org/10.1007/978-981-13-3459-7_3).
- Kareem Ahmed, Tao Li, Thy Ton, Quan Guo, Kai-Wei Chang, Parisa Kordjamshidi, Vivek Srikumar, Guy Van den Broeck, and Sameer Singh. Pylon: A pytorch framework for learning with constraints. In Douwe Kiela, Marco Ciccone, and Barbara Caputo, editors, *Proceedings of the NeurIPS 2021 Competitions and Demonstrations Track*, volume 176 of *Proceedings of Machine Learning Research*, pages 319–324. PMLR, 06–14 Dec 2022. URL <https://proceedings.mlr.press/v176/ahmed22a.html>.
- Luis M. Augusto. From symbols to knowledge systems: A. newell and h. a. Simon’s contribution to symbolic AI. *Journal of Knowledge Structures and Systems*, 2(1):29–62, 2021.
- Samy Badreddine, Artur d’Avila Garcez, Luciano Serafini, and Michael Spranger. Logic tensor networks. *Artificial Intelligence*, 303:103649, 2022. ISSN 0004-3702. doi: <https://doi.org/10.1016/j.artint.2021.103649>. URL <https://www.sciencedirect.com/science/article/pii/S0004370221002009>.
- Emily M. Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, FAccT ’21, page 610–623, New York, NY, USA, 2021. Association for Computing Machinery. ISBN 9781450383097. doi: 10.1145/3442188.3445922. URL <https://doi.org/10.1145/3442188.3445922>.

- Bikram Pratim Bhuyan, Amar Ramdane-Cherif, Ravi Tomar, and T. P. Singh. Neuro-symbolic artificial intelligence: a survey. *Neural Computing and Applications*, 36(21): 12809–12844, July 2024. ISSN 1433-3058. doi: 10.1007/s00521-024-09960-z. URL <https://doi.org/10.1007/s00521-024-09960-z>.
- Grady Booch, Francesco Fabiano, Lior Horesh, Kiran Kate, Jonathan Lenchner, Nick Linck, Andreas Loreggia, Keerthiram Murgesan, Nicholas Mattei, Francesca Rossi, and Biplav Srivastava. Thinking fast and slow in ai. *Proceedings of the AAAI Conference on Artificial Intelligence*, 35(17):15042–15046, May 2021. doi: 10.1609/aaai.v35i17.17765. URL <https://ojs.aaai.org/index.php/AAAI/article/view/17765>.
- Djallel Bouneffouf and Charu C. Aggarwal. Survey on applications of neurosymbolic artificial intelligence, 2022. URL <https://arxiv.org/abs/2209.12618>.
- Ivan Bratko and Stephen Muggleton. Applications of inductive logic programming. *Commun. ACM*, 38(11):65–70, November 1995. ISSN 0001-0782. doi: 10.1145/219717.219771. URL <https://doi.org/10.1145/219717.219771>.
- E. Burattini, A. de Francesco, and M. De Gregorio. Nsl: a neuro-symbolic language for monotonic and non-monotonic logical inferences. In *VII Brazilian Symposium on Neural Networks, 2002. SBRN 2002. Proceedings.*, pages 256–261, 2002. doi: 10.1109/SBRN.2002.1181487.
- William F Clocksin and Christopher S Mellish. *Programming in PROLOG*. Springer Science & Business Media, 2003.
- William W Cohen, Fan Yang, and Kathryn Rivard Mazaitis. Tensorlog: Deep learning meets probabilistic dbs. *arXiv preprint arXiv:1707.05390*, 2017.
- Andrew Cropper and Sebastijan Dumančić. Inductive logic programming at 30: A new introduction. *J. Artif. Int. Res.*, 74, September 2022. ISSN 1076-9757. doi: 10.1613/jair.1.13507. URL <https://doi.org/10.1613/jair.1.13507>.
- Adnan Darwiche. Sdd: A new canonical representation of propositional knowledge bases. In *IJCAI Proceedings-International Joint Conference on Artificial Intelligence*, volume 22, page 819, 2011.
- Luc De Raedt, Angelika Kimmig, and Hannu Toivonen. Problog: a probabilistic prolog and its application in link discovery. In *Proceedings of the 20th International Joint Conference on Artificial Intelligence, IJCAI’07*, page 2468–2473, San Francisco, CA, USA, 2007. Morgan Kaufmann Publishers Inc.
- Jason Eisner. Parameter estimation for probabilistic finite-state transducers. In *Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics*, pages 1–8, 2002.
- Francesco Fabiano, Vishal Pallagani, Marianna Bergamaschi Ganapini, Lior Horesh, Andrea Loreggia, Keerthiram Murgesan, Francesca Rossi, and Biplav Srivastava. Plan-SOFAI: A neuro-symbolic planning architecture. In *Neuro-Symbolic Learning and Reasoning in*

- the era of Large Language Models*, 2023. URL <https://openreview.net/forum?id=ORAhayOH4x>.
- Hossein Rajaby Faghihi, Aliakbar Nafar, Chen Zheng, Roshanak Mirzaee, Yue Zhang, Andrzej Uszok, Alexander Wan, Tanawan Premisri, Dan Roth, and Parisa Kordjamshidi. Gluecons: A generic benchmark for learning under constraints, 2023. URL <https://arxiv.org/abs/2302.10914>.
- Hossein Rajaby Faghihi, Aliakbar Nafar, Andrzej Uszok, Hamid Karimian, and Parisa Kordjamshidi. Prompt2demodel: Declarative neuro-symbolic modeling with natural language, 2024. URL <https://arxiv.org/abs/2407.20513>.
- Chuyu Fang and Song-Chun Yu. Large language models are neurosymbolic reasoners. *arXiv preprint arXiv:2401.09334*, 2024. URL <https://arxiv.org/html/2401.09334v1>.
- Michael Frazier and Leonard Pitt. Learning from entailment: An application to propositional horn sentences. In *Proceedings of the Tenth International Conference on International Conference on Machine Learning*, pages 120–127, 1993.
- Eleonora Giunchiglia, Alex Tatomir, Mihaela Cătălina Stoian, and Thomas Lukasiewicz. Ccn+: A neuro-symbolic framework for deep learning with requirements. *International Journal of Approximate Reasoning*, 171:109124, 2024. ISSN 0888-613X. doi: <https://doi.org/10.1016/j.ijar.2024.109124>. URL <https://www.sciencedirect.com/science/article/pii/S0888613X24000112>. Synergies between Machine Learning and Reasoning.
- Todd J. Green, Grigoris Karvounarakis, and Val Tannen. Provenance semirings. In *Proceedings of the Twenty-Sixth ACM SIGMOD-SIGACT-SIGART Symposium on Principles of Database Systems*, PODS '07, page 31–40, New York, NY, USA, 2007. Association for Computing Machinery. ISBN 9781595936851. doi: 10.1145/1265530.1265535. URL <https://doi.org/10.1145/1265530.1265535>.
- Quan Guo, Hossein Rajaby Faghihi, Yue Zhang, Andrzej Uszok, and Parisa Kordjamshidi. Inference-masked loss for deep structured output learning. In Christian Bessiere, editor, *Proceedings of the Twenty-Ninth International Joint Conference on Artificial Intelligence, IJCAI-20*, pages 2754–2761. International Joint Conferences on Artificial Intelligence Organization, 7 2020. doi: 10.24963/ijcai.2020/382. URL <https://doi.org/10.24963/ijcai.2020/382>. Main track.
- Gurobi Optimization, LLC. Gurobi Optimizer Reference Manual, 2024. URL <https://www.gurobi.com>.
- Frederick Hayes-Roth. Rule-based systems. *Communications of the ACM*, 28(9):921–932, 1985.
- Pascal Hitzler and Md Kamruzzaman Sarker. Neuro-symbolic artificial intelligence: The state of the art. 2022.
- Delower Hossain and Jake Y Chen. A study on neuro-symbolic artificial intelligence: Healthcare perspectives, 2025. URL <https://arxiv.org/abs/2503.18213>.

- Joy Hsu, Jiayuan Mao, Joshua B. Tenenbaum, and Jiajun Wu. What’s left? concept grounding with logic-enhanced foundation models, 2023. URL <https://arxiv.org/abs/2310.16035>.
- Jiani Huang, Ziyang Li, Binghong Chen, Karan Samel, Mayur Naik, Le Song, and Xujie Si. Scallop: From probabilistic deductive databases to scalable differentiable reasoning. In M. Ranzato, A. Beygelzimer, Y. Dauphin, P.S. Liang, and J. Wortman Vaughan, editors, *Advances in Neural Information Processing Systems*, volume 34, pages 25134–25145. Curran Associates, Inc., 2021. URL [https://proceedings.neurips.cc/paper\\_files/paper/2021/file/d367eef13f90793bd8121e2f675f0dc2-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2021/file/d367eef13f90793bd8121e2f675f0dc2-Paper.pdf).
- Adam Ishay, Zhun Yang, and Joohyung Lee. Leveraging large language models to generate answer set programs. In Pierre Marquis, Tran Cao Son, and Gabriele Kern-Isberner, editors, *Proceedings of the 20th International Conference on Principles of Knowledge Representation and Reasoning, KR 2023*, Proceedings of the International Conference on Knowledge Representation and Reasoning, pages 374–383. Association for the Advancement of Artificial Intelligence, 2023. doi: 10.24963/kr.2023/37.
- Prashani Jayasingha, Bogdan Iancu, and Johan Lilius. Neurosymbolic approaches in ai design – an overview. In *2025 IEEE Symposium on Trustworthy, Explainable and Responsible Computational Intelligence (CITREx Companion)*, pages 1–5, 2025. doi: 10.1109/CITRExCompanion65208.2025.10981497.
- Justin Johnson, Bharath Hariharan, Laurens van der Maaten, Li Fei-Fei, C. Lawrence Zitnick, and Ross Girshick. Clevr: A diagnostic dataset for compositional language and elementary visual reasoning. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, July 2017.
- Daniel Kahneman. *Thinking, fast and slow*. macmillan, 2011.
- Danial Kamali, Elham J. Barezi, and Parisa Kordjamshidi. Nesycoco: A neuro-symbolic concept composer for compositional generalization. *Proceedings of the AAAI Conference on Artificial Intelligence*, 39(4):4184–4193, Apr. 2025. doi: 10.1609/aaai.v39i4.32439. URL <https://ojs.aaai.org/index.php/AAAI/article/view/32439>.
- Henry Kautz. The third ai summer: Aaai robert s. engelmore memorial lecture. *AI Magazine*, 43(1):105–125, Mar. 2022. doi: 10.1002/aaai.12036. URL <https://ojs.aaai.org/aimagazine/index.php/aimagazine/article/view/19122>.
- Angelika Kimmig, Guy Van den Broeck, and Luc De Raedt. An algebraic prolog for reasoning about possible worlds. *Proceedings of the AAAI Conference on Artificial Intelligence*, 25(1):209–214, Aug. 2011. doi: 10.1609/aaai.v25i1.7852. URL <https://ojs.aaai.org/index.php/AAAI/article/view/7852>.
- Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete Xiao, Spencer Whitehead, Alexander C. Berg, Wan-Yen Lo, Piotr Dollár, and Ross Girshick. Segment anything, 2023. URL <https://arxiv.org/abs/2304.02643>.



- Phokion G. Kolaitis and Moshe Y. Vardi. On the expressive power of datalog: tools and a case study. In *Proceedings of the Ninth ACM SIGACT-SIGMOD-SIGART Symposium on Principles of Database Systems*, PODS '90, page 61–71, New York, NY, USA, 1990. Association for Computing Machinery. ISBN 0897913523. doi: 10.1145/298514.298542. URL <https://doi.org/10.1145/298514.298542>.
- Parisa Kordjamshidi, Dan Roth, and Hao Wu. Saul: Towards declarative learning based programming. volume 2015, 12 2015.
- Parisa Kordjamshidi, Daniel Khashabi, Christos Christodoulopoulos, Bhargav Mangipudi, Sameer Singh, and Dan Roth. Better call Saul: Flexible programming for learning and inference in NLP. In Yuji Matsumoto and Rashmi Prasad, editors, *Proceedings of COLING 2016, the 26th International Conference on Computational Linguistics: Technical Papers*, pages 3030–3040, Osaka, Japan, December 2016. The COLING 2016 Organizing Committee. URL <https://aclanthology.org/C16-1285/>.
- Parisa Kordjamshidi, Dan Roth, and Kristian Kersting. Declarative learning-based programming as an interface to ai systems. *Frontiers in Artificial Intelligence*, 5, 2019. URL <https://api.semanticscholar.org/CorpusID:195069123>.
- Parisa Kordjamshidi, Dan Roth, and Kristian Kersting. Declarative learning-based programming as an interface to ai systems. *Frontiers in Artificial Intelligence*, Volume 5 - 2022, 2022. ISSN 2624-8212. doi: 10.3389/frai.2022.755361. URL <https://www.frontiersin.org/journals/artificial-intelligence/articles/10.3389/frai.2022.755361>.
- Luis C. Lamb, Artur Garcez, Marco Gori, Marcelo Prates, Pedro Avelar, and Moshe Vardi. Graph neural networks meet neural-symbolic computing: A survey and perspective, 2021. URL <https://arxiv.org/abs/2003.00330>.
- Guillaume Lample and François Charton. Deep learning for symbolic mathematics. In *International Conference on Learning Representations*, 2020. URL <https://openreview.net/forum?id=S1eZYeHFDS>.
- Y. Lecun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998. doi: 10.1109/5.726791.
- Yann LeCun, Yoshua Bengio, and Geoffrey Hinton. Deep learning. *Nature*, 521(7553):436–444, 2015. doi: 10.1038/nature14539. URL <https://doi.org/10.1038/nature14539>.
- Bo Li, Peng Qi, Bo Liu, Shuai Di, Jingen Liu, Jiquan Pei, Jinfeng Yi, and Bowen Zhou. Trustworthy ai: From principles to practices. *ACM Comput. Surv.*, 55(9), January 2023a. ISSN 0360-0300. doi: 10.1145/3555803. URL <https://doi.org/10.1145/3555803>.
- Ziyang Li, Jiani Huang, and Mayur Naik. Scallop: A language for neurosymbolic programming. *Proc. ACM Program. Lang.*, 7(PLDI), June 2023b. doi: 10.1145/3591280. URL <https://doi.org/10.1145/3591280>.

- Ziyang Li, Jiani Huang, Jason Liu, Felix Zhu, Eric Zhao, William Dodds, Neelay Velingker, Rajeev Alur, and Mayur Naik. Relational programming with foundational models. *Proceedings of the AAAI Conference on Artificial Intelligence*, 38(9):10635–10644, March 2024. ISSN 2159-5399. doi: 10.1609/aaai.v38i9.28934. URL <http://dx.doi.org/10.1609/aaai.v38i9.28934>.
- Priscila Lima, Mariela Morveli-Espinoza, Glaucia K E Pereira, and Felipe França. Satyrus: A sat-based neuro-symbolic architecture for constraint processing. pages 137–142, 01 2005. doi: 10.1109/ICHIS.2005.97.
- Runtao Liu, Chenxi Liu, Yutong Bai, and Alan L. Yuille. Clevr-ref+: Diagnosing visual reasoning with referring expressions. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2019.
- Zhen Lu, Imran Afridi, Hong Jin Kang, Ivan Ruchkin, and Xi Zheng. Surveying neuro-symbolic approaches for reliable artificial intelligence of things. *Journal of Reliable Intelligent Environments*, 10(3):257–279, 2024. doi: 10.1007/s40860-024-00231-1. URL <https://doi.org/10.1007/s40860-024-00231-1>.
- Robin Manhaeve, Sebastijan Dumančić, Angelika Kimmig, Thomas Demeester, and Luc De Raedt. Neural probabilistic logic programming in deepproblog. *Artificial Intelligence*, 298:103504, 2021. ISSN 0004-3702. doi: <https://doi.org/10.1016/j.artint.2021.103504>. URL <https://www.sciencedirect.com/science/article/pii/S0004370221000552>.
- Jiayuan Mao, Chuang Gan, Pushmeet Kohli, Joshua B. Tenenbaum, and Jiajun Wu. The neuro-symbolic concept learner: Interpreting scenes, words, and sentences from natural supervision, 2019. URL <https://arxiv.org/abs/1904.12584>.
- Matthias Minderer, Alexey Gritsenko, Austin Stone, Maxim Neumann, Dirk Weissenborn, Alexey Dosovitskiy, Aravindh Mahendran, Anurag Arnab, Mostafa Dehghani, Zhuoran Shen, Xiao Wang, Xiaohua Zhai, Thomas Kipf, and Neil Houlsby. Simple open-vocabulary object detection with vision transformers, 2022. URL <https://arxiv.org/abs/2205.06230>.
- Roshanak Mirzaee and Parisa Kordjamshidi. Disentangling extraction and reasoning in multi-hop spatial reasoning. In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 3379–3397, Singapore, December 2023. Association for Computational Linguistics. doi: 10.18653/v1/2023.findings-emnlp.221. URL <https://aclanthology.org/2023.findings-emnlp.221/>.
- Yatin Nandwani, Abhishek Pathak, Mausam, and Parag Singla. A primal dual formulation for deep learning with constraints. In H. Wallach, H. Larochelle, A. Beygelzimer, F. d'Alché-Buc, E. Fox, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc., 2019. URL [https://proceedings.neurips.cc/paper\\_files/paper/2019/file/cf708fc1decf0337aded484f8f4519ae-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2019/file/cf708fc1decf0337aded484f8f4519ae-Paper.pdf).

- Allen Newell. Physical symbol systems. *Cognitive Science*, 4(2):135–183, 1980. ISSN 0364-0213. doi: [https://doi.org/10.1016/S0364-0213\(80\)80015-2](https://doi.org/10.1016/S0364-0213(80)80015-2). URL <https://www.sciencedirect.com/science/article/pii/S0364021380800152>.
- Raymond Ng and V.S. Subrahmanian. Probabilistic logic programming. *Information and Computation*, 101(2):150–201, 1992. ISSN 0890-5401. doi: [https://doi.org/10.1016/0890-5401\(92\)90061-J](https://doi.org/10.1016/0890-5401(92)90061-J). URL <https://www.sciencedirect.com/science/article/pii/089054019290061J>.
- Shan-Hwei Nienhuys-Cheng and Roland de Wolf. *What is inductive logic programming?* Springer, 1997.
- OpenAI, Josh Achiam, and Steven Adler et al. Gpt-4 technical report, 2024. URL <https://arxiv.org/abs/2303.08774>.
- Liangming Pan, Alon Albalak, Xinyi Wang, and William Wang. Logic-LM: Empowering large language models with symbolic solvers for faithful logical reasoning. In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 3806–3824, Singapore, December 2023a. Association for Computational Linguistics. doi: 10.18653/v1/2023.findings-emnlp.248. URL <https://aclanthology.org/2023.findings-emnlp.248/>.
- Liangming Pan, Alon Albalak, Xinyi Wang, and William Yang Wang. Logic-lm: Empowering large language models with symbolic solvers for faithful logical reasoning, 2023b. URL <https://arxiv.org/abs/2305.12295>.
- Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Köpf, Edward Yang, Zach DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. Pytorch: An imperative style, high-performance deep learning library, 2019. URL <https://arxiv.org/abs/1912.01703>.
- Fabio Petroni, Tim Rocktäschel, Patrick Lewis, Anton Bakhtin, Yuxiang Wu, Alexander H. Miller, and Sebastian Riedel. Language models as knowledge bases?, 2019. URL <https://arxiv.org/abs/1909.01066>.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision, 2021. URL <https://arxiv.org/abs/2103.00020>.
- Hossein Rajaby Faghihi, Quan Guo, Andrzej Uszok, Aliakbar Nafar, and Parisa Kordjamshidi. DomiKnowS: A library for integration of symbolic domain knowledge in deep learning. In Heike Adel and Shuming Shi, editors, *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 231–241, Online and Punta Cana, Dominican Republic, November 2021. Association for Computational Linguistics. doi: 10.18653/v1/2021.emnlp-demo.27. URL <https://aclanthology.org/2021.emnlp-demo.27/>.

- Saratha Sathasivam. Learning rules comparison in neuro-symbolic integration. *International Journal of Applied Physics and Mathematics*, 1(2):129, 2011.
- Luciano Serafini and Artur S d’Avila Garcez. Learning and reasoning with logic tensor networks. In *Conference of the italian association for artificial intelligence*, pages 334–348. Springer, 2016.
- Amir Shpilka, Amir Yehudayoff, et al. Arithmetic circuits: A survey of recent results and open questions. *Foundations and Trends® in Theoretical Computer Science*, 5(3–4): 207–388, 2010.
- David Silver, Aja Huang, Chris J. Maddison, Arthur Guez, Laurent Sifre, George van den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, Sander Dieleman, Dominik Grewe, John Nham, Nal Kalchbrenner, Ilya Sutskever, Timothy Lillicrap, Madeleine Leach, Koray Kavukcuoglu, Thore Graepel, and Demis Hassabis. Mastering the game of go with deep neural networks and tree search. *Nature*, 529(7587):484–489, 2016. doi: 10.1038/nature16961. URL <https://doi.org/10.1038/nature16961>.
- Paul Smolensky, Moontae Lee, Xiaodong He, Wen-tau Yih, Jianfeng Gao, and li Deng. Basic reasoning with tensor product representations. 01 2016. doi: 10.48550/arXiv.1601.02745.
- Jiankai Sun, Hao Sun, Tian Han, and Bolei Zhou. Neuro-symbolic program search for autonomous driving decision module design. In Jens Kober, Fabio Ramos, and Claire Tomlin, editors, *Proceedings of the 2020 Conference on Robot Learning*, volume 155 of *Proceedings of Machine Learning Research*, pages 21–30. PMLR, 16–18 Nov 2021. URL <https://proceedings.mlr.press/v155/sun21a.html>.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, Dan Bikel, Lukas Blecher, Cristian Canton Ferrer, Moya Chen, Guillem Cucurull, David Esiobu, Jude Fernandes, Jeremy Fu, Wenyin Fu, Brian Fuller, Cynthia Gao, Vedanuj Goswami, Naman Goyal, Anthony Hartshorn, Saghar Hosseini, Rui Hou, Hakan Inan, Marcin Kardas, Viktor Kerkez, Madian Khabsa, Isabel Kloumann, Artem Korenev, Punit Singh Koura, Marie-Anne Lachaux, Thibaut Lavril, Jenya Lee, Diana Liskovich, Yinghai Lu, Yuning Mao, Xavier Martinet, Todor Mihaylov, Pushkar Mishra, Igor Molybog, Yixin Nie, Andrew Poulton, Jeremy Reizenstein, Rashi Rungta, Kalyan Saladi, Alan Schelten, Ruan Silva, Eric Michael Smith, Ranjan Subramanian, Xiaoqing Ellen Tan, Binh Tang, Ross Taylor, Adina Williams, Jian Xiang Kuan, Puxin Xu, Zheng Yan, Iliyan Zarov, Yuchen Zhang, Angela Fan, Melanie Kambadur, Sharan Narang, Aurelien Rodriguez, Robert Stojnic, Sergey Edunov, and Thomas Scialom. Llama 2: Open foundation and fine-tuned chat models, 2023. URL <https://arxiv.org/abs/2307.09288>.
- Pascal Van Hentenryck, Helmut Simonis, and Mehmet Dincbas. Constraint satisfaction using constraint logic programming. *Artificial Intelligence*, 58(1):113–159, 1992. ISSN 0004-3702. doi: [https://doi.org/10.1016/0004-3702\(92\)90006-J](https://doi.org/10.1016/0004-3702(92)90006-J). URL <https://www.sciencedirect.com/science/article/pii/000437029290006J>.

- Joseph Weizenbaum. Eliza—a computer program for the study of natural language communication between man and machine. *Commun. ACM*, 9(1):36–45, January 1966. ISSN 0001-0782. doi: 10.1145/365153.365168. URL <https://doi.org/10.1145/365153.365168>.
- Lionel Wong, Gabriel Grand, Alexander K. Lew, Noah D. Goodman, Vikash K. Mansinghka, Jacob Andreas, and Joshua B. Tenenbaum. From word models to world models: Translating from natural language to the probabilistic language of thought, 2023. URL <https://arxiv.org/abs/2306.12672>.
- Fangzhi Xu, Zhiyong Wu, Qiushi Sun, Siyu Ren, Fei Yuan, Shuai Yuan, Qika Lin, Yu Qiao, and Jun Liu. Symbol-LLM: Towards foundational symbol-centric interface for large language models. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 13091–13116, Bangkok, Thailand, August 2024a. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.707. URL <https://aclanthology.org/2024.acl-long.707/>.
- Jingyi Xu, Zilu Zhang, Tal Friedman, Yitao Liang, and Guy Van den Broeck. A semantic loss function for deep learning with symbolic knowledge. In Jennifer Dy and Andreas Krause, editors, *Proceedings of the 35th International Conference on Machine Learning*, volume 80 of *Proceedings of Machine Learning Research*, pages 5502–5511. PMLR, 10–15 Jul 2018. URL <https://proceedings.mlr.press/v80/xu18h.html>.
- Xinyu Xu, Pan Wang, Shusen Dong, Ningyu Wu, Yue Zhang, and Baobao Chang. Faithful logical reasoning via symbolic chain-of-thought. *arXiv preprint arXiv:2405.18357*, 2024b. URL <https://arxiv.org/abs/2405.18357>.
- Xiaocheng Yang, Bingsen Chen, and Yik-Cheung Tam. Arithmetic reasoning with LLM: Prolog generation & permutation. In Kevin Duh, Helena Gomez, and Steven Bethard, editors, *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 2: Short Papers)*, pages 699–710, Mexico City, Mexico, June 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.naacl-short.61. URL <https://aclanthology.org/2024.naacl-short.61/>.
- Liang Yao, Jiazhen Peng, Chengsheng Mao, and Yuan Luo. Exploring large language models for knowledge graph completion, 2025. URL <https://arxiv.org/abs/2308.13916>.
- Kexin Yi, Jiajun Wu, Chuang Gan, Antonio Torralba, Pushmeet Kohli, and Joshua B. Tenenbaum. Neural-symbolic vqa: Disentangling reasoning from vision and language understanding. In *Advances in Neural Information Processing Systems*, pages 1039–1050, 2018.
- Hanlin Zhang, Jiani Huang, Ziyang Li, Mayur Naik, and Eric Xing. Improved logical reasoning of language models via differentiable symbolic programming, 2023. URL <https://arxiv.org/abs/2305.03742>.

Danna Zheng, Mirella Lapata, and Jeff Z. Pan. How reliable are llms as knowledge bases? re-thinking facutality and consistency, 2024. URL <https://arxiv.org/abs/2407.13578>.

## Appendix A. Example Task

NeSy frameworks formulate problems differently based on their implementation and setup. In **DomiKnowS**, the problem is reformulated as a logical constraint solving problem. This is done by representing the domain as a graph  $G(V, E)$ , where the nodes are the concepts in the domain and the edges are the relationships between them. Each node can have properties. The final logical constraint formulation is done using the defined concepts. In **DeepProbLog**, the problem is viewed as a combination of perception and reasoning, where the perception is the neural component that is fed as neural predicates into the reasoning component made of probabilistic logic programming with ProbLog. To solve a problem in DeepProbLog, the problem needs to be conceptualized as a separation of the neural and logical reasoning components. In **Scallop**, similar to DeepProbLog, the problem is viewed as a combination of the neural and the symbolic component. In **LEFT**, the problem is limited to the visual question answering domain. Here, the neural model is composed of feature extractors, a classifier for objects and relations into concepts, and a first-order logic program generator when given the question. In this section, we will take a look at how the problem formulation looks like in each of these frameworks for a common task. **Note** that we exclude LEFT for this due to the domain-specific nature of the framework.

### A.1. MNIST Sum

The MNIST Sum task is an extension of the classic MNIST handwritten digit recognition task (Lecun et al., 1998) where on being given two images of digits, the task is to output their sum that is a whole number. The training example consists of the two images of the digits and the ground-truth label of their sum. The individual labels of the digits are not available for training.

#### A.1.1. DOMIKNOWS

**Problem Specification.** DomiKnowS formulates the problem using graph representations of concepts, relations, and logic. For performing the MNIST Sum task in DomiKnowS, the first concept defined is *image* concept representing visual information. The *digit* concept, a subclass of image, is introduced to represent the output class, ranging from 0 to 9. To establish relationships between digit images, the *image pair* concept is defined as an edge connecting two digit concepts. The sum concept is then introduced under image pair to represent the summation of the two digit concepts and the ground-truth output of the program. For this task, three constraints are defined. The first two constraints utilize *exactL* to ensure that the predicted digit and sum values belong to only one valid class. Another constraint enforces that the expected sum value matches the sum of the two digit predictions. This is implemented using *ifL* constraints, which verify whether the predicted digits form one of the possible solutions for a valid sum. If multiple solutions exist, the *orL* constraint ensures that at least one of the answers corresponds to the predicted digits. Code for defining the graph concept and constraints can be found in Listing 1.



```

with Graph(name='global') as graph:
    image_batch = Concept(name='image_batch')
    image = Concept(name='image')

    image_contains, = image_batch.contains(image)

# digit classes 0-9
    digit = image(name='digits',
                  ConceptClass=EnumConcept,
                  values=digits)

    image_pair = Concept(name='pair')
    pair_d0, pair_d1 = image_pair.has_a(digit0=image, digit1=
        image)

# sum value classes 0-18
    s = image_pair(name='summations',
                   ConceptClass=EnumConcept,
                   values=summations)

    exactL(*[digit.__getattr__(d) for d in digits])
    exactL(*[s.__getattr__(d) for d in summations])

#fixedL(s)
    FIXED = True
    fixedL(s("x", eqL(image_pair, "summationEquality", {True})
        ), active = FIXED)

    for sum_val in range(config.summationRange):
        sum_combinations = []

        sum_nm = summations[sum_val]

        for d0_val in range(sum_val + 1):
            d1_val = sum_val - d0_val

            if d0_val >= len(digits) or d1_val >= len(digits):
                continue

            d0_nm = digits[d0_val]
            d1_nm = digits[d1_val]

            # for each combination of digits that sum to
            sum_val add constraint to list

```

```

sum_combinations.append(andL(getattr(digit , d0_nm)
    (path=('x' , pair_d0)),
                                getattr(digit , d1_nm)
                                (path=('x' ,
                                    pair_d1))
                                ))

print(sum_val , '-' , sum_combinations)

# if the given summation value is some value , then the
digits must be one of a set of
# digit pairs that add to that value
# i.e. if sum val = s, d0 = 0 and d1 = s or d0 = 1 and
      d1 = s-1 ...
# e.g. if sum val = 1, d0 = 0 and d1 = 1 or d0 = 1 and
      d0 = 0
if len(sum_combinations) == 1:
    ifL(
        getattr(s , sum_nm)('x') ,
        sum_combinations[0]
    )
else:
    ifL(
        getattr(s , sum_nm)('x') ,
        orL(*sum_combinations)
    )

```

Listing 1: Python Code for full graph of MNIST sum implemented in DomiKnowS including logical constraints

**Neural Modeling.** The model declaration comprises standard neural modeling components, including data loading, pre-processing, neural network definition, and loss function specification. The process begins with the *ReaderSensor*, which reads the input image. Next, a relation concept is defined using another sensor, *JointSensor*, to establish connections between images. The module learner is then employed to generate an initial prediction for the digit concept, which is subsequently passed to another sensor, *FunctionalSensor*, to compute the sum of two images. The associated code is provided in Listing 2.

```

class Net(torch.nn.Module):
    def __init__(self):
        super().__init__()

        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)

```

```

self.pool = nn.MaxPool2d(2, 2)

self.lin1 = nn.Linear(256, 128)
self.lin2 = nn.Linear(128, 10)

self.relu = nn.ReLU()

self.drop = nn.Dropout(p=0.2)

self.norm = nn.LayerNorm(256)

def forward(self, x):
    x = torch.squeeze(x, dim=0)

    x = x.reshape(2, 1, 28, 28)

    x = self.conv1(x)
    x = self.relu(x)
    x = self.pool(x)

    x = self.conv2(x)
    x = self.relu(x)
    x = self.pool(x)

    x = x.reshape(2, -1)

    x = self.norm(x)

    x = self.lin1(x)
    x = self.relu(x)

    x = self.drop(x)

    y_digit = self.lin2(x)

    return y_digit

class SumLayer(torch.nn.Module):
    def __init__(self):
        super().__init__()

        self.lin1 = nn.Linear(20, 64)
        self.lin2 = nn.Linear(64, 19)

```

```

self.relu = nn.ReLU()

def forward(self, digits, do_time=True):
    digit0 = torch.unsqueeze(digits[0, :], dim=0)
    digit1 = torch.unsqueeze(digits[1, :], dim=0)

    x = torch.cat((digit0, digit1), dim=1)

    x = self.lin1(x)
    x = self.relu(x)

    y_sum = self.lin2(x)

    #return torch.zeros((1, 19), requires_grad=True)
    return y_sum

class SumLayerExplicit(torch.nn.Module):
    def __init__(self, device='cpu'):
        super().__init__()
        self.device = device

    def forward(self, digits, do_time=True):
        digit0 = torch.unsqueeze(digits[0, :], dim=0)
        digit1 = torch.unsqueeze(digits[1, :], dim=0)

        digit0 = F.softmax(digit0, dim=1)
        digit1 = F.softmax(digit1, dim=1)

        digit0 = torch.reshape(digit0, (10, 1))
        digit1 = torch.reshape(digit1, (1, 10))
        d = torch.matmul(digit0, digit1)
        d = d.repeat(1, 1, 1, 1)
        f = torch.flip(torch.eye(10), dims=(0,)).repeat(1, 1,
            1, 1)
        conv_diag_sums = F.conv2d(d, f.to(self.device),
            padding=(9, 0), groups=1)[..., 0]

        out = torch.squeeze(conv_diag_sums, dim=0)
        return out

class NBSoftCrossEntropyLoss(NBCrossEntropyLoss):
    def __init__(self, prior_weight=1.0, *args, **kwargs):
        super().__init__(*args, **kwargs)

```

```

        self.prior_weight = prior_weight

    def forward(self, input, target, *args, **kwargs):
        if target.dim() == 1:
            return super().forward(input, target, *args, **
                                    kwargs)

        epsilon = 1e-5
        input = input.view(-1, input.shape[-1])
        input = input.clamp(min=epsilon, max=1-epsilon)

        logprobs = F.log_softmax(input, dim=1)
        return self.prior_weight * -(target * logprobs).sum()
            / input.shape[0]

class NBSoftCrossEntropyIMLoss(BCEWithLogitsIMLoss):
    def __init__(self, prior_weight=1.0, *args, **kwargs):
        super().__init__(*args, **kwargs)

        self.prior_weight = prior_weight

    def forward(self, input, inference, target, weight=None):
        if target.dim() == 1:
            num_classes = input.shape[-1]
            target = target.to(dtype=torch.long)
            target = F.one_hot(target, num_classes=num_classes)

            return super().forward(input, inference, target,
                                    weight=weight)

        return super().forward(input, inference, target,
                                weight=weight) * self.prior_weight

def print_and_output(x, f=lambda x: x.shape, do_print=False):
    if do_print:
        print(prefix + str(f(x)))
    return x

def build_program(sum_setting=None, digit_labels=False, device
                  ='cpu', use_fixedL=True, test=False):

```

```

image['pixels'] = ReaderSensor(keyword='pixels')

def make_batch(pixel):
    return pixel.flatten().unsqueeze(0), torch.ones((1,
        len(pixel)))
image_batch['pixels', image_contains.reversed] =
    JointSensor(image['pixels'], forward=make_batch)

image['logits'] = ModuleLearner('pixels', module=Net())

def make_pairs(*inputs):
    return torch.tensor([[True, False]]), torch.tensor([[
        False, True]])

image_pair[pair_d0.reversed, pair_d1.reversed] =
    JointSensor(image['pixels'], forward=make_pairs)

image_pair['summation_label'] = ReaderSensor(keyword='
    summation')

image['digit_label'] = ReaderSensor(keyword='digit')

image[digit] = FunctionalSensor('logits', forward=lambda x
    : x)

if digit_labels:
    image[digit] = FunctionalSensor('digit_label', forward
        =lambda x: x, label=True)

if use_fixedL and test:
    # during test time, set model output to be the
    summation label
    def manual_fixedL(s):
        res = torch.zeros((1, 19))
        res[0, s] = 1
        return res

    image_pair[s] = FunctionalSensor('summation_label',
        forward=manual_fixedL)
else:
    if sum_setting == 'explicit':
        image_pair[s] = ModuleLearner(image['logits'],
            module=SumLayerExplicit(device=device))
    elif sum_setting == 'baseline':

```



```

        image_pair[s] = ModuleLearner(image['logits'],
                                      module=SumLayer())
    else:
        image_pair[s] = FunctionalSensor(forward=lambda:
                                         torch.ones(1, config.summationRange)) # dummy
                                         values to populate

    if use_fixedL:
        image_pair[s] = ReaderSensor(keyword='summation',
                                      label=True)
        image_pair['summationEquality'] = FunctionalSensor(
            forward=lambda: torch.ones(1,1))

    return graph, image, image_pair, image_batch

```

Listing 2: MNIST Sum code for DomiKnowS framework to run this task

#### A.1.2. DEEPPROBLOG

**Problem Specification.** DeepProbLog formulates a problem regarding probabilistic facts, neural facts, and neural annotated disjunctions (nAD). In the MNIST Sum task, the fact  $X$  is defined to represent the input image. A neural network function is then introduced to map  $X$  to its corresponding digit, denoted as  $\text{digit}(X, Y)$ . To enforce constraints about the summation and the ground-truth sum, a function is defined to compute the sum of two digits. Code for this part is shown in Listing 3.

```

nn(m_digit, [X], Y, [0.....9]) :: digit(X,Y).
addition(X,Y,Z) :- digit(X,X2), digit(Y,Y2), Z is X2+Y2.

```

Listing 3: Facts and Rules in DeepProbLog

**Neural Modeling.** The neural modeling follows a standard neural network setup, such as a CNN-based classifier. It is preceded by data loading and pre-processing, which are performed separately from the ProbLog program. Thus, the neural model used in DeepProbLog can be initialized independently of the DeepProbLog model. Once the neural model is initialized, the framework passes it along with a probabilistic program as input. The probabilistic program consists of facts and rules, similar to the code in Listing 3. Details of the DeepProbLog modeling code can be found in Listing 4.

```

class Model(object):
    def __init__(
        self,
        program_string: Union[str, os.PathLike],
        networks: Collection[Network],
        embeddings: Optional[TermEmbedder] = None,
        load: bool = True,
    ):
        """

```

```

:param program_string: A string representing a
    DeepProbLog program or the path to a file
    containing a program.
:param networks: A collection of networks that will be
    used to evaluate the neural predicates.
:param embeddings: A TermEmbedder used to embed Terms
    in the program.
:param load: If true, then it will attempt to load the
    program from 'program_string',
    else, it will consider program_string to be the
    program itself.
"""
self.networks = dict()
if load:
    self.program: LogicProgram = PrologFile(str(
        program_string))
else:
    self.program: LogicProgram = PrologString(
        program_string)
self.parameters = []
self.parameter_groups = []
self._extract_parameters()
for network in networks:
    self.networks[network.name] = network
    network.model = self
self.solver: Optional[Solver] = None
self.eval_mode = False
self.embeddings = embeddings
self.tensor_sources = dict()
self.optimizer = Optimizer(self)

def get_embedding(self, term: Term):
    return self.embeddings.get_embedding(term)

def evaluate_nn(self, to_evaluate: List[Tuple[Term, Term
]]) :
    """
    :param to_evaluate: List of neural predicates to
        evaluate
    :return: A dictionary with the elements of to_evaluate
        as keys, and the output of the NN as values.
    """
    result = dict()
    evaluations = defaultdict(list)

```

```

# Group inputs per net to send in batch
for net_name, inputs in to_evaluate:
    net = self.networks[str(net_name)]
    if net.det:
        tensor_name = Term("nn", net_name, inputs)
        if tensor_name not in self.solver.engine.
            tensor_store:
                evaluations[net_name].append(inputs)
    else:
        if inputs in net.cache:
            result[(net_name, inputs)] = net.cache[
                inputs]
            del net.cache[inputs]
        else:
            evaluations[net_name].append(inputs)
for net in evaluations:
    network = self.networks[str(net)]
    out = network([term2list(x, False) for x in
        evaluations[net]])
    for i, k in enumerate(evaluations[net]):
        if network.det:
            tensor_name = Term("nn", net, k)
            self.solver.engine.tensor_store.store(out[
                i], tensor_name)
        else:
            result[(net, k)] = out[i]
return result

def set_engine(self, engine: Engine, **kwargs):
    """
    Initializes the solver of this model with the given
    engine and additional arguments.
    :param engine: The engine that will be used to ground
        queries in this model.
    :param kwargs: Additional arguments passed to the
        solver.
    :return:
    """
    self.solver = Solver(self, engine, **kwargs)
    register_tensor_predicates(engine)

def solve(self, batch: Sequence[Query]) -> List[Result]:
    return self.solver.solve(batch)

def ground_dataset(self, dataset: Dataset):

```

```

total_time = 0
compile_times = []
ground_times = []
for q in dataset.to_queries():
    start = time.time()
    result = self.solver.cache.get(q)
    total_time += time.time() - start
    if not result.from_cache:
        compile_times.append(result.compile_time)
        ground_times.append(result.ground_time)
return {
    "total_time": total_time,
    "ground_times": ground_times,
    "compile_times": compile_times,
}

def save_state(self, filename: Union[str, PathLike, IO[
bytes]], complete=False):
    """
    Saves the state of this model to a zip file with the
    given filename. This only includes the
    probabilistic
    parameters and all parameters of the neural
    networks, but not the model architecture or
    neural architectures
    :param filename: The filename to save the model to.
    :param complete: If true, save neural networks with
        information needed to resume training.
    :return:
    """
    check_path(filename)
    with ZipFile(filename, "w") as zipf:
        with zipf.open("parameters", "w") as f:
            pickle.dump(self.parameters, f)
        for n in self.networks:
            with zipf.open(n, "w") as f:
                self.networks[n].save(f, complete=complete
                )

def load_state(self, filename: Union[str, PathLike, IO[
bytes]]):
    """
    Restore the state of this model from the given
    filename. This only includes the probabilistic
    parameters

```

```

        and all parameters of the neural networks, but not
        the model architecture or neural architectures

:param filename: The filename to restore the model
from.
:return:
"""
with ZipFile(filename) as zipf:
    with zipf.open("parameters") as f:
        self.parameters = pickle.load(f)
    for n in self.networks:
        with zipf.open(n) as f:
            self.networks[n].load(BytesIO(f.read()))

def eval(self):
    """
    Set the mode of all networks in the model to eval.
    """
    self.eval_mode = True
    for n in self.networks:
        self.networks[n].eval()
    self.solver.engine.eval()

def train(self):
    """
    Set the mode of all networks in the model to train.
    :return:
    """
    self.eval_mode = False
    for n in self.networks:
        self.networks[n].train()
    self.solver.engine.train()

def register_foreign(
    self, func: Callable, function_name: str, arity_in:
        int, arity_out: int
):
    self.solver.engine.register_foreign(func,
        function_name, arity_in, arity_out)

def __str__(self):
    return "\n".join(str(line) for line in self.program)

def get_tensor(self, term: Term) -> torch.Tensor:
    """

```

```

        :param term: A term of the form tensor(_).
        If the tensor is of the form tensor(a(*args)), then it
            well look into tensor source a.
        :return: Returns the stored tensor identifier by the
            term.
        """
        if len(term.args) > 0 and term.args[0].functor in self
            .tensor_sources:
            return self.tensor_sources[term.args[0].functor][
                term.args[0].args]
        return self.solver.get_tensor(term)

def store_tensor(self, tensor: torch.Tensor) -> Term:
    """
    Stores a tensor in the tensor store and returns and
        identifier.
    :param tensor: The tensor to store.
    :return: The Term that is the identifier by which this
        tensor can be uniquely identified in the logic.
    """
    return Term("tensor", Constant(self.solver.engine.
        tensor_store.store(tensor)))

def add_tensor_source(
    self, name: str, source: Union[ImageDataset, Mapping[
        Any, torch.Tensor]]
):
    """
    Adds a named tensor source to the model.
    :param name: The name of the added tensor source.
    :param source: The tensor source to add
    :return:
    """
    self.tensor_sources[name] = source

def get_hyperparameters(self) -> dict:
    """
    Recursively build a dictionary containing the most
        important hyperparameters in the model.
    :return: A dictionary that contains the values of the
        most important hyperparameters of the model.
    """
    parameters = dict()
    parameters["solver"] = (

```



```

        None if self.solver is None else self.solver.
            get_hyperparameters()
    )
    parameters["networks"] = [
        self.networks[network].get_hyperparameters() for
            network in self.networks
    ]
    parameters["program"] = self.program.to_prolog()
    return parameters

def hyperparameters_to_file(self, filename):
    """
    Write the output of the get_hyperparameter() method in
    JSON format to a file.
    :param filename: The path to write the hyperparameters
        to.
    :return:
    """
    with open(filename, "w") as f:
        f.write(json.dumps(self.get_hyperparameters()))

def _extract_parameters(self):
    translated = SimpleProgram()
    for n in self.program:
        if type(n) is Term:
            if (
                n.probability is not None
                and type(n.probability) is Term
                and n.probability.functor == "t"
            ):
                i = self._add_parameter(n.probability.args
                    [0])
                p = n.probability.with_args(Constant(i))
                n = n.with_probability(p)
                translated.add_statement(n)
            elif type(n) is Clause:
                if (
                    n.head.probability is not None
                    and type(n.head.probability) is Term
                    and n.head.probability.functor == "t"
                ):
                    i = self._add_parameter(n.head.probability
                        .args[0])
                    p = n.head.probability.with_args(Constant(
                        i))

```

```

        head = n.head.with_probability(p)
        n = Clause(head, n.body)
        translated.add_statement(n)
    elif type(n) is Or:
        new_list = []
        new_group = []
        for x in n.to_list():
            if (
                x.probability is not None
                and type(x.probability) is Term
                and x.probability.functor == "t"
            ):
                i = self._add_parameter(x.probability.
                    args[0])
                new_group.append(i)
                p = x.probability.with_args(Constant(i))
                new_list.append(x.with_probability(p))
            else:
                new_list.append(x)
        if len(new_group) > 0:
            self.parameter_groups.append(new_group)
        n = Or.from_list(new_list)
        translated.add_statement(n)
    else:
        translated.add_statement(n)
    self.program = translated

def _add_parameter(self, val: Constant):
    i = len(self.parameters)
    try:
        val = float(val)
    except InstantiationException:
        val = random()
    self.parameters.append(val)
    return i

```

Listing 4: Example of code for neural model in DeepProbLog for MNIST Sum task.

### A.1.3. SCALLOP

**Problem Specification.** Scallop formulates the problem in terms of relations, values, and (Horn) rules derived from Datalog. As discussed earlier, the concepts and constraints defined in this framework are similar to those in DeepProbLog. However, these rules can be directly embedded into a Scallop program through its API. The process begins by establishing the concepts *digit1* and *digit2* to represent the digit values of two given images. Constraints

are then defined based on the summation of these two values, which must be equal to the *sum\_2* concept, serving as the ground truth for this task. The code in Listing 5 provides a portion of the implementation for defining these concepts and constraints.

```
self.scl_ctx.add_relation('`digit_1`', int, input_mapping=list
    (range(10)))
self.scl_ctx.add_relation('`digit_2`', int, input_mapping=list
    (range(10)))
self.scl_ctx.add_rule('`sum_2(a + b) :- digit_1(a), digit_2(b)
    ``')
self.sum_2 = self.scl_ctx.forward_function('`sum_2`',
    output_mapping=[(i,) for i in range(19)])
```

Listing 5: Code for defining the concept and constraints in Scallop framework for MNIST sum task

**Neural Modeling.** Unlike DeepProbLog, the neural modeling is integrated with Scallop’s relation and rule declaration. The neural modeling remains a standard neural network. Details of the modeling code can be found in Listing 6.

```
mnist_img_transform = torchvision.transforms.Compose([
    torchvision.transforms.ToTensor(),
    torchvision.transforms.Normalize(
        (0.1307,), (0.3081,)
    )
])

class MNISTSum2Dataset(torch.utils.data.Dataset):
    def __init__(
        self,
        root: str,
        train: bool = True,
        transform: Optional[Callable] = None,
        target_transform: Optional[Callable] = None,
        download: bool = False,
    ):
        # Contains a MNIST dataset
        self.mnist_dataset = torchvision.datasets.MNIST(
            root,
            train=train,
            transform=transform,
            target_transform=target_transform,
            download=download,
        )
        self.index_map = list(range(len(self.mnist_dataset)))
        random.shuffle(self.index_map)
```

```

def __len__(self):
    return int(len(self.mnist_dataset) / 2)

def __getitem__(self, idx):
    # Get two data points
    (a_img, a_digit) = self.mnist_dataset[self.index_map[idx *
        2]]
    (b_img, b_digit) = self.mnist_dataset[self.index_map[idx *
        2 + 1]]

    # Each data has two images and the GT is the sum of two
    # digits
    return (a_img, b_img, a_digit + b_digit)

@staticmethod
def collate_fn(batch):
    a_imgs = torch.stack([item[0] for item in batch])
    b_imgs = torch.stack([item[1] for item in batch])
    digits = torch.stack([torch.tensor(item[2]).long() for
        item in batch])
    return ((a_imgs, b_imgs), digits)

def mnist_sum_2_loader(data_dir, batch_size_train,
    batch_size_test):
    train_loader = torch.utils.data.DataLoader(
        MNISTSum2Dataset(
            data_dir,
            train=True,
            download=True,
            transform=mnist_img_transform,
        ),
        collate_fn=MNISTSum2Dataset.collate_fn,
        batch_size=batch_size_train,
        shuffle=True
    )

    test_loader = torch.utils.data.DataLoader(
        MNISTSum2Dataset(
            data_dir,
            train=False,
            download=True,
            transform=mnist_img_transform,
        ),

```

```

        collate_fn=MNISTSum2Dataset.collate_fn ,
        batch_size=batch_size_test ,
        shuffle=True
    )

    return train_loader , test_loader

class MNISTNet(nn.Module):
    def __init__(self):
        super(MNISTNet, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, kernel_size=5)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=5)
        self.fc1 = nn.Linear(1024, 1024)
        self.fc2 = nn.Linear(1024, 10)

    def forward(self, x):
        x = F.max_pool2d(self.conv1(x), 2)
        x = F.max_pool2d(self.conv2(x), 2)
        x = x.view(-1, 1024)
        x = F.relu(self.fc1(x))
        x = F.dropout(x, p = 0.5, training=self.training)
        x = self.fc2(x)
        return F.softmax(x, dim=1)

class MNISTSum2Net(nn.Module):
    def __init__(self, provenance, k):
        super(MNISTSum2Net, self).__init__()

        # MNIST Digit Recognition Network
        self.mnist_net = MNISTNet()

        # Scallop Context
        self.scl_ctx = scallopy.ScallopContext(provenance=
            provenance, k=k)
        self.scl_ctx.add_relation("digit_1", int, input_mapping=
            list(range(10)))
        self.scl_ctx.add_relation("digit_2", int, input_mapping=
            list(range(10)))
        self.scl_ctx.add_rule("sum_2(a+-b) :- - digit_1(a), - digit_2
            (b)")

        # The `sum_2` logical reasoning module

```

```

self.sum_2 = self.scl_ctx.forward_function("sum_2",
    output_mapping=[(i,) for i in range(19)], jit=args.jit,
    dispatch=args.dispatch)

def forward(self, x: Tuple[torch.Tensor, torch.Tensor]):
    (a_imgs, b_imgs) = x

    # First recognize the two digits
    a_distrs = self.mnist_net(a_imgs) # Tensor 64 x 10
    b_distrs = self.mnist_net(b_imgs) # Tensor 64 x 10

    # Then execute the reasoning module; the result is a size
    19 tensor
    return self.sum_2(digit_1=a_distrs, digit_2=b_distrs) #
        Tensor 64 x 19

def bce_loss(output, ground_truth):
    (_, dim) = output.shape
    gt = torch.stack([torch.tensor([1.0 if i == t else 0.0 for i
        in range(dim)]) for t in ground_truth])
    return F.binary_cross_entropy(output, gt)

def nll_loss(output, ground_truth):
    return F.nll_loss(output, ground_truth)

class Trainer():
    def __init__(self, train_loader, test_loader, model_dir,
        learning_rate, loss, k, provenance):
        self.model_dir = model_dir
        self.network = MNISTSum2Net(provenance, k)
        self.optimizer = optim.Adam(self.network.parameters(), lr=
            learning_rate)
        self.train_loader = train_loader
        self.test_loader = test_loader
        self.best_loss = 10000000000
        if loss == "nll":
            self.loss = nll_loss
        elif loss == "bce":
            self.loss = bce_loss
        else:
            raise Exception(f"Unknown loss function - `{loss}`")

```

```

def train_epoch(self, epoch):
    self.network.train()
    iter = tqdm(self.train_loader, total=len(self.train_loader))
    for (data, target) in iter:
        self.optimizer.zero_grad()
        output = self.network(data)
        loss = self.loss(output, target)
        loss.backward()
        self.optimizer.step()
        iter.set_description(f"[Train-Epoch-{epoch}] - Loss: {loss.item():.4f}")

def test_epoch(self, epoch):
    self.network.eval()
    num_items = len(self.test_loader.dataset)
    test_loss = 0
    correct = 0
    with torch.no_grad():
        iter = tqdm(self.test_loader, total=len(self.test_loader))
        for (data, target) in iter:
            output = self.network(data)
            test_loss += self.loss(output, target).item()
            pred = output.data.max(1, keepdim=True)[1]
            correct += pred.eq(target.data.view_as(pred)).sum()
            perc = 100. * correct / num_items
            iter.set_description(f"[Test-Epoch-{epoch}] - Total-loss: {test_loss:.4f}, - Accuracy: {correct}/{num_items} ({perc:.2f}%)")
        if test_loss < self.best_loss:
            self.best_loss = test_loss
            torch.save(self.network, os.path.join(model_dir, "sum_2_best.pt"))

def train(self, n_epochs):
    self.test_epoch(0)
    for epoch in range(1, n_epochs + 1):
        self.train_epoch(epoch)
        self.test_epoch(epoch)
    
```

Listing 6: Example of code for neural model in Scallop for MNIST Sum task.